



WEB CRAWLER

- Team members: Yu Tzu Li, Mayukh Sinha
- Course: CSYE7200
- Tech Stack: Scala, Apache Pekko, Apache Kafka
- Date: December 2025

OLD ARCHITECTURE OVERVIEW



ARCHITECTURE OVERVIEW



DIFFERENCE

Aspect	NEW	OLD
Master functionality	deduplication + frontier + metrics + worker registry	Only publishes URLs
Kafka topics	tasks / results / registration / heartbeats	urls / results
Worker capabilities	Fetching, parsing, metrics, heartbeat, registration	Fetching, parsing
Fault tolerance	Worker heartbeat → can detect failures	No worker health monitoring
Deduplication	Bloom Filter for deduplication	No deduplication (may fetch duplicates)
Crawler engine loop	Yes (extracted links return to frontier)	No (results do not feed back into new URLs)

MASTER NODE

- The Master Node contains three key components:
 - Master Actor
 - URL Frontier + Bloom Filter
 - Metrics Collector

MASTER ACTOR

- It is responsible for:
 - Receiving worker registrations
 - Receiving worker heartbeats
 - Receiving crawl results from workers
 - Sending FetchTask messages to Kafka
 - Managing the URL Frontier

URL FRONTIER + BLOOM FILTER

- **Frontier:** A queue that stores URLs waiting to be crawled
- **Bloom Filter:** Used to check whether a URL has already been visited → deduplication

METRICS COLLECTOR

- tracks worker states
- counts the total number of processed tasks
- monitors overall crawling performance

KAFKA

- tasks (Master → Workers)
- registration (Workers → Master)
- results (Workers → Master)
- heartbeats (Workers → Master)

TASKS (MASTER → WORKERS)

- Master Actor takes URLs popped from the Frontier and publishes them to the tasks topic.
- Workers consume tasks from this topic.

REGISTRATION (WORKERS → MASTER)

- When a worker starts, it sends registration information such as, workerId

RESULTS (WORKERS → MASTER)

- After a worker finishes fetching a page, it sends:
 - the URL
 - the fetched content
 - extracted outlinks
 - status code
 - latency
- When the Master consumes results:
 - it updates metrics
 - it extracts new URLs and pushes them into the Frontier (after Bloom Filter deduplication)

HEARTBEATS (WORKERS → MASTER)

- Workers periodically send heartbeats.
- The Master uses them to monitor worker health

WORKERS

- Each worker (Worker 1, Worker 2, ... Worker N) has its own complete processing pipeline.
- The Worker Node contains three key components:
 - Worker Actor
 - Http Fetcher
 - Page Parser
 - Local Metrics

WORKER ACTOR

- consume tasks → call HttpFetcher
- call PageParser → produce results (links / content)
- publish results to Kafka
- send heartbeats periodically
- send registration on startup

HTTP FETCHER

- The component that actually issues HTTP requests.
- Flow:
 - Worker Actor receives a URL
 - Http Fetcher performs a GET request
 - The response body is returned to PageParser

PAGE PARSER

- Responsible for:
- parsing HTML
- extracting outlinks
- packaging results as CrawlResult
- The Worker Actor then publishes the parsed result to Kafka's results topic.

LOCAL METRICS

- These metrics can be sent to Kafka
 - HTTP latency
 - success/failure count
 - status code

INTERNET

- This represents the external websites that workers send HTTP requests to.
- Diagram flow:
 - Worker → HttpFetcher → Internet → response → PageParser

BENCHMARK ENVIRONEMNT

Component	Details
Processor	Apple M3 Pro, 11-core CPU (6 performance + 5 efficiency cores)
Graphics	Integrated Apple 14-core GPU
Memory	18 GB Unified Memory
Storage	512 GB SSD
Architecture	ARM-based Apple Silicon (aarch64)

ACCEPTANCE CRITERIA (20 WORKERS)

Metric	Value
Workers	20
Throughput (pages/sec)	83.47
Average Latency (ms)	2946
Total URLs Crawled	1,579
Elapsed Time (sec)	18

ACCEPTANCE CRITERIA (DETAILED STATS)

Category	Count
URLs Crawled	1,579
└ Succeeded	1,116
└ Failed	463
Success Rate	70.7%
Bytes Downloaded	258.18 MB
Links Extracted	205,622

ACCEPTANCE CRITERIA (TOP DOMAINS)

Domain	Count
developer.mozilla.org	49
nodejs.org	50
reactjs.org	50
blog.cloudflare.com	48
news.ycombinator.com	144
<u>www.elastic.co</u>	50
arstechnica.com	50
redis.io	49
en.wikipedia.org	65
github.com	81

ACCEPTANCE CRITERIA (30 WORKERS)

Metric	Value
Workers	30
Throughput (pages/sec)	134.81
Average Latency (ms)	4690
Total URLs Crawled	3,580
Elapsed Time (sec)	26

ACCEPTANCE CRITERIA (DETAILED STATS)

Category	Count
URLs Crawled	3,580
└ Succeeded	2,343
└ Failed	1,237
Success Rate	65.4%
Bytes Downloaded	569.35 MB
Links Extracted	374,637

ACCEPTANCE CRITERIA (TOP DOMAINS)

Domain	Count
www.linkedin.com	61
www.postgresql.org	200
nodejs.org	51
spark.apache.org	65
www.reddit.com	185
www.coursera.org	56
news.ycombinator.com	145
docs.python.org	61
en.wikipedia.org	221
github.com	98

ACCEPTANCE CRITERIA

Metric	Planned Acceptance Criteria	Actual (20 Workers)	Actual (30 Workers)	Observation / Justification
Throughput (pages/sec)	300–600 pages/sec with 10 workers (expected to scale proportionally)	83.47	134.81	Real-world pages had heavier content, network delays, and external rate limits. Throughput still scaled upward, showing the architecture works under realistic conditions.
Average Latency (ms)	< 900 ms at sustained load	2,946 ms	4,690 ms	Latency targets were based on synthetic assumptions; real websites responded slowly. Increased concurrency introduced natural queuing delays. System remained stable throughout.
Speed-Up Ratio	≥ 3× speed-up when scaling 5 → 20 workers	Sub-linear	Sub-linear (20→30 = ~1.6×)	External bottlenecks (DNS time, server rate limits, bandwidth) limit speed-up, not our system. Internal architecture scales correctly until external limits are hit.

MILESTONES

Week	Original Planned Milestone	Revised Milestone (Updated After Realistic Performance Findings)
Week 1	Set up Kafka and Pekko environments. Define basic Master-Worker actor communication.	Same as planned. Environment + communication layer initialized successfully.
Week 2	Implement crawling logic and integrate HTML parsing.	Same as planned. Basic crawling pipeline built and validated with small test loads.
Week 3	Add result collection and metrics monitoring.	Major Architecture Update: <ul style="list-style-type: none">Identified bottlenecks during early load testing.Migrated to more efficient batching & back-pressure model.Introduced async I/O optimization, connection pooling, and improved rate-limiting strategies.Designed a more realistic, scalable crawling pipeline to handle real-world latency and response variability.

MILESTONES

Week

Original Planned Milestone

Week 4

Run scalability tests and prepare final report + demo.

End Goal


Distributed web crawler suitable for real-world use (news aggregation, indexing, etc.)

Revised Milestone (Updated After Realistic Performance Findings)

- Execute revised scalability tests using updated architecture.
- Capture latency, throughput, and scaling metrics under realistic workloads.
- Finalize performance analysis, comparisons to acceptance criteria, and prepare demo.

Robust, fault-tolerant distributed crawler capable of handling real-world network variability, with improved performance stability and scalability.

UNIT TESTS (CRAWLER-API)



A screenshot of a terminal window displaying the results of a unit test run. The window has a dark theme with light-colored text. At the top, there is a toolbar with various icons. Below the toolbar, the title "Test Results" is shown, followed by a detailed list of test cases and their execution times.

Test Case	Time
parse() should extract title, description, text, links, images, headings, metaTags	49 ms
extractLinks should resolve absolute + relative links, remove invalid, skip extensions	42 ms
extractImages should extract absolute URLs	2 ms
extractHeadings should collect h1, h2, h3 text only	1 ms
extractText should remove script/style/noscript content	0 ms
getMetaDescription should return description content when available	1 ms
extractMetaTags should extract both name and property attributes	0 ms
extractLinks should filter invalid protocols but may include baseUrl via Jsoup behavior	1 ms
normalizeUrl should lowercase scheme/host but preserve original path case	1 ms

On the right side of the terminal window, there is a status bar with the message "9 tests passed 9 tests" and the path "/Library/Java/JavaVirtualMachines". Below that, it says "Testing started at 1:25 PM".

UNIT TESTS (CRAWLER-CORE)

Test Results		
✓	PrioritizedUrlSpec	202 ms
✓	Higher priority comes first	7 ms
✓	Tie priority → earlier timestamp first	0 ms
✓	UrlFrontierSpec	41 ms
✓	addUrl should add valid URL and dedupe duplicates	35 ms
✓	maxDepth enforcement	0 ms
✓	normalizeUrl should lowercase + remove trailing slash	0 ms
✓	getNext returns URLs in priority order	1 ms
✓	getBatch returns up to requested size	5 ms
✓	CrawlerMetricsSpec	8 ms
✓	recordSuccess should increment counters	7 ms
✓	recordFailure should increment failure counters	1 ms
✓	getAvgFetchTimeMs computation	0 ms
✓	MetricsReporterSpec	3 ms
✓	MetricsReporter start/stop should not throw	3 ms
✓	MetricsSnapshotSpec	0 ms
✓	MetricsSnapshot stores values correctly	0 ms
✓	ControlMessagesSpec	2 ms
✓	ShutdownWorker messageType	1 ms
✓	PauseWorker messageType	0 ms
✓	ResumeWorker messageType	0 ms
✓	UrlBatch messageType	1 ms
✓	FetchResultSpec	1 ms
✓	FetchResult should expose messageType	1 ms
✓	FetchResult should store fields correctly	0 ms
✓	FetchTaskSpec	1 ms
✓	FetchTask should expose correct messageType	1 ms
✓	FetchTask should maintain all given fields	0 ms

UNIT TESTS (CRAWLER-CORE)

▶	✓ FetchResult should store fields correctly	0 ms
▶	✓ FetchTaskSpec	1ms
▶	✓ FetchTask should expose correct messageType	1ms
▶	✓ FetchTask should maintain all given fields	0 ms
▶	✓ MessageCodecsSpec	136 ms
▶	✓ Round-trip encode/decode FetchTask	134 ms
▶	✓ Round-trip encode/decode FetchResult	1ms
▶	✓ PauseWorker custom encoder/decoder	0 ms
▶	✓ Unknown _type should return error	1ms
▶	✓ RegisterWorkerSpec	1ms
▶	✓ RegisterWorker should expose messageType	1ms
▶	✓ RegisterWorker should store metadata	0 ms
▶	✓ WorkerHeartbeatSpec	1ms
▶	✓ WorkerHeartbeat should expose messageType	1ms
▶	✓ WorkerHeartbeat should store data correctly	0 ms
▶	✓ WorkerRegisteredSpec	1ms
▶	✓ WorkerRegistered should expose messageType	0 ms
▶	✓ WorkerRegistered should store fields	1ms

UNIT TESTS (CRAWLER-MASTER)

A screenshot of a Java IDE (IntelliJ IDEA) displaying unit test results. The interface includes a toolbar at the top with various icons, a left sidebar with navigation buttons, and a main window showing a tree view of test results.

The test results are as follows:

Test Suite	Test Case	Duration	
MasterCommandSpec	WorkerRegistration wraps RegisterWorker	7 ms	
	HeartbeatReceived wraps WorkerHeartbeat	7 ms	
	ResultReceived wraps FetchResult	0 ms	
	WorkerInfoSpec	WorkerInfo should store fields	0 ms
		WorkerInfoSpec	1 ms
		WorkerRegistrySpec	1 ms
	WorkerRegistrySpec	register should create worker entry	59 ms
		assignTask should add task to assignedTasks	4 ms
		completeTask should remove assigned task	0 ms
		getAvailableWorker should return worker if below concurrency	0 ms
getAvailableWorker should return None when busy		0 ms	
getDeadWorkerTasks returns tasks assigned to dead workers		1 ms	
getAvailableWorker		54 ms	

Summary: 10 tests passed, 10 tests total, 67 ms.

Testing started at 1:18 PM ...

UNIT TESTS (CRAWLER-WORKER)

The screenshot shows a dark-themed IDE interface with a "Run" tab selected, displaying the results of ScalaTests in the 'test' package. There are 3 test suites and 8 individual test cases listed, all of which have passed. The total execution time for all tests is 219 ms.

Test Suite / Case	Status	Duration
WorkerCommandSpec	✓ Passed	13 ms
ProcessTask stores FetchTask	✓ Passed	12 ms
FetchComplete stores task and response	✓ Passed	1 ms
SendHeartbeat is a singleton	✓ Passed	0 ms
FetchResponseSpec	✓ Passed	2 ms
FetchResponse stores all fields correctly in success case	✓ Passed	1 ms
FetchResponse handles error and missing body/contentType	✓ Passed	0 ms
isSuccess returns true only for status codes 200–299	✓ Passed	1 ms
HttpFetcherSpec	✓ Passed	204 ms
FakeFetcher returns FetchResponse	✓ Passed	203 ms
FakeFetcher close() should set flag	✓ Passed	1 ms

Additional UI elements visible include a toolbar with icons for running, stopping, and navigating tests, and a status bar at the bottom right indicating "/Library/Java/Testing started".